

# Analyze A/B Test Results

## Table of Contents

- [Introduction](#)
- [Part I - Probability](#)
- [Part II - A/B Test](#)
- [Part III - Regression](#)

### Introduction

This is an A/B test to determine if using a new website will help convert the users from free users to paying members. 3 different concepts will be used to reach a conclusion if the new website design is more effective than the original design at converting users to paid members.

#### Part I - Probability

To get started, let's import our libraries.

```
In [3]: import pandas as pd
import numpy as np
import random
import matplotlib.pyplot as plt
%matplotlib inline
# we are setting the seed to assure you get the same answers on quizzes as we set up
random.seed(42)
```

1. Now, read in the `ab_data.csv` data. Store it in `df`. Use your dataframe to answer the questions in Quiz 1 of the classroom.

a. Read in the dataset and take a look at the top few rows here:

```
In [4]: df = pd.read_csv('ab_data.csv')
df.head()
```

```
Out[4]:
```

|   | user_id | timestamp                  | group     | landing_page | converted |
|---|---------|----------------------------|-----------|--------------|-----------|
| 0 | 851104  | 2017-01-21 22:11:48.556739 | control   | old_page     | 0         |
| 1 | 804228  | 2017-01-12 08:01:45.159739 | control   | old_page     | 0         |
| 2 | 661590  | 2017-01-11 16:55:06.154213 | treatment | new_page     | 0         |
| 3 | 853541  | 2017-01-08 18:28:03.143765 | treatment | new_page     | 0         |
| 4 | 864975  | 2017-01-21 01:52:26.210827 | control   | old_page     | 1         |

b. Use the cell below to find the number of rows in the dataset.

```
In [5]: df.shape
Out[5]: (294478, 5)
```

c. The number of unique users in the dataset.

```
In [6]: df.nunique()
Out[6]:
```

|   | user_id | timestamp | group | landing_page | converted |
|---|---------|-----------|-------|--------------|-----------|
| 0 | 290584  | 294478    | 2     | 2            | 2         |
| 1 | 2       | 2         | 2     | 2            | 2         |
| 2 | int64   | int64     | int64 | int64        | int64     |

d. The proportion of users converted.

```
In [7]: df['converted'].mean()
Out[7]: 0.1196591935605512
```

e. The number of times the `new_page` and `treatment` don't match.

```
In [8]: # create a dataframe where control and new page match. Then check how many rows are 1
df1 = df.query('group == "control" and landing_page == "new_page"')
df1.count()
```

```
Out[8]:
```

|   | user_id   | timestamp | group     | landing_page | converted |
|---|-----------|-----------|-----------|--------------|-----------|
| 0 | 1928      | 1928      | 1928      | 1928         | 1928      |
| 1 | converted | converted | converted | converted    | converted |
| 2 | int64     | int64     | int64     | int64        | int64     |

```
In [9]: # create a dataframe where treatment and old page match. Then check how many rows are 1
df2 = df.query('group == "treatment" and landing_page == "old_page"')
df2.count()
```

```
Out[9]:
```

|   | user_id   | timestamp | group     | landing_page | converted |
|---|-----------|-----------|-----------|--------------|-----------|
| 0 | 1965      | 1965      | 1965      | 1965         | 1965      |
| 1 | group     | group     | group     | group        | group     |
| 2 | converted | converted | converted | converted    | converted |
| 3 | int64     | int64     | int64     | int64        | int64     |

f. Do any of the rows have missing values?

```
In [10]: df.isnull().sum()
Out[10]:
```

|   | user_id   | timestamp | group     | landing_page | converted |
|---|-----------|-----------|-----------|--------------|-----------|
| 0 | 0         | 0         | 0         | 0            | 0         |
| 1 | group     | group     | group     | group        | group     |
| 2 | converted | converted | converted | converted    | converted |
| 3 | int64     | int64     | int64     | int64        | int64     |

2. For the rows where `treatment` does not match with `new_page` or `control` does not match with `old_page`, we cannot be sure if this row truly received the new or old page. Use Quiz 2 in the classroom to figure out how we should handle these rows.

a. Now use the answer to the quiz to create a new dataset that meets the specifications from the quiz. Store your new dataframe in `df2`.

```
In [11]: # drop the dataframe where control and new page match from the original dataframe and
df1 = df.drop(df1.index)
df1.shape
```

```
Out[11]: (292550, 5)
```

```
In [12]: # now from this newly created df1 drop the dataframe where treatment and old page match
df2 = df1.drop(df2.index)
df2.shape
```

```
Out[12]: (290585, 5)
```

```
In [13]: # Double Check all of the correct rows were removed - this should be 0
df2[(df2['group'] == 'treatment') && (df2['landing_page'] == 'new_page')] == False).sum()
Out[13]: 0
```

3. Use `df2` and the cells below to answer questions for Quiz3 in the classroom.

a. How many unique `user_ids` are in `df2`?

```
In [14]: df2.nunique()
Out[14]:
```

|   | user_id      | timestamp    | group        | landing_page | converted    |
|---|--------------|--------------|--------------|--------------|--------------|
| 0 | 290584       | 290585       | 2            | 2            | 2            |
| 1 | landing_page | landing_page | landing_page | landing_page | landing_page |
| 2 | converted    | converted    | converted    | converted    | converted    |
| 3 | int64        | int64        | int64        | int64        | int64        |

b. There is one `user_id` repeated in `df2`. What is it?

```
In [15]: df2['user_id'].duplicated().sum()
Out[15]: 1
```

c. What is the row information for the repeat `user_id`?

```
In [16]: duplicate = df2[df2.duplicated(['user_id'])]
duplicate
```

```
Out[16]:
```

|      | user_id | timestamp                  | group     | landing_page | converted |
|------|---------|----------------------------|-----------|--------------|-----------|
| 2893 | 773192  | 2017-01-14 02:55:59.590927 | treatment | new_page     | 0         |

```
In [17]: # find the shape of df2
df2.shape
Out[17]: (290585, 5)
```

d. Remove one of the rows with a duplicate `user_id`, but keep your dataframe as `df2`.

```
In [18]: df2 = df2.drop(duplicate.index)
In [19]: # check the shape again to make sure the duplicated row has been removed
df2.shape
Out[19]: (290584, 5)
```

4. Use `df2` in the cells below to answer the quiz questions related to Quiz 4 in the classroom.

a. What is the probability of an individual converting regardless of the page they receive?

```
In [20]: df2['converted'].mean()
Out[20]: 0.11959708724499628
```

b. Given that an individual was in the `control` group, what is the probability they converted?

```
In [21]: df2.query('group=="control"')['converted'].mean()
Out[21]: 0.1203863045004612
```

c. Given that an individual was in the `treatment` group, what is the probability they converted?

```
In [22]: df2.query('group=="treatment"')['converted'].mean()
Out[22]: 0.1188080655150564
```

d. What is the probability that an individual received the new new page?

```
In [23]: new = df2.query('landing_page == "new_page"')
len(new)/len(df2)
Out[23]: 0.500061944222688
```

e. Consider your results from parts (a) through (d) above, and explain below whether you think there is sufficient evidence to conclude that the new treatment page leads to more conversions. 0.1203863045004612 - 0.1188080655150564 = 0.001578238985355567 This is less than a 1 % change in conversion rate. In fact the control group seems to have a better conversion rate than the new treatment page. so far there is no evidence that the new treatment page leads to more conversions

### Part II - A/B Test

Notice that because of the time stamp associated with each event, you could technically run a hypothesis test continuously as each observation was observed.

However, then the hard question is do you stop as soon as one page is considered significantly better than another or does it need to happen consistently for a certain amount of time? How long do you run to render a decision that neither page is better than another?

These questions are the difficult parts associated with A/B tests in general.

1. For now, consider you need to make the decision just based on all the data provided. If you want to assume that the old page is better unless the new page proves to be definitely better at a Type I error rate of 5%, what should your null and alternative hypotheses be? You can state your hypothesis in terms of words or in terms of  $p_{old}$  and  $p_{new}$ , which are the converted rates for the old and new pages.

Null:  $p_{old} \geq p_{new}$

Alternate:  $p_{old} < p_{new}$

2. Assume under the null hypothesis,  $p_{new}$  and  $p_{old}$  both have a "true" success rates equal to the converted to success rate regardless of page - that is  $p_{new}$  and  $p_{old}$  are equal. Furthermore, assume they are equal to the conversion rate in `ab_data.csv` regardless of the page.

Use a sample size for each page equal to the ones in `ab_data.csv`.

Perform the sampling distribution for the difference in converted between the two pages over 10,000 iterations of calculating an estimate from the null.

Use the cells below to provide the necessary parts of this simulation. If this doesn't make complete sense right now, don't worry - you are going to work through the problems below to complete this problem. You can use Quiz 5 in the classroom to make sure you are on the right track.

a. What is the conversion rate for  $p_{new}$  under the null?

```
In [24]: p_new = df2['converted'].mean()
p_new
```

```
Out[24]: 0.11959708724499628
```

b. What is the conversion rate for  $p_{old}$  under the null?

```
In [25]: p_old = p_new
p_old
```

```
Out[25]: 0.11959708724499628
```

c. What is  $n_{new}$ , the number of individuals in the treatment group?

```
In [26]: # user_id is arbitrarily selected as one of the columns used for count. any of the ot
# would have given the same result
n_new = df2.query('group == "treatment"')['user_id'].count()
n_new
```

```
Out[26]: 145310
```

d. What is  $n_{old}$ , the number of individuals in the control group?

```
In [27]: # user_id is arbitrarily selected as one of the columns used for count. any of the ot
# would have given the same result
n_old = df2.query('group == "control"')['user_id'].count()
n_old
```

```
Out[27]: 145274
```

e. Simulate  $n_{new}$  transactions with a conversion rate of  $p_{new}$  under the null. Store these  $n_{new}$  1's and 0's in `new_page_converted`.

```
In [28]: new_page_converted = np.random.choice([0,1], size=n_new, p = [p_new, p_new])
new_page_converted
```

```
Out[28]: array([0, 0, 1, ..., 0, 0, 0])
```

f. Simulate  $n_{old}$  transactions with a conversion rate of  $p_{old}$  under the null. Store these  $n_{old}$  1's and 0's in `old_page_converted`.

```
In [29]: old_page_converted = np.random.choice([0,1], size=n_old, p = [p_old, p_old])
old_page_converted
```

```
Out[29]: array([0, 0, 0, ..., 0, 0, 0])
```

g. Find  $p_{new} - p_{old}$  for your simulated values from part (e) and (f).

```
In [30]: p_new - p_old
Out[30]: 0.0
```


h. Create 10,000  $p_{new} - p_{old}$  values using the same simulation process you used in parts (a) through (g) above. Store all  $p_{new} - p_{old}$  values in a NumPy array called `p_diffs`.

```
In [31]: # this is commented out and replaced with code below. this is an alternative method to
# the binomial method is preferred as it avoids using a python loop which runs much s
# diff = []
# for i in range(10000):
#     new_page_converted = np.random.choice([0,1], size=n_new, p = [p_new, p_new]).me
#     old_page_converted = np.random.choice([0,1], size=n_old, p = [p_old, p_old]).me
#     diff = new_page_converted - old_page_converted
#     p_diffs.append(diff)
```

```
In [32]: new_converted_simulation = np.random.binomial(n_new, p_new, 10000)/n_new
old_converted_simulation = np.random.binomial(n_old, p_old, 10000)/n_old
p_diffs = new_converted_simulation - old_converted_simulation
```

i. Plot a histogram of the `p_diffs`. Does this plot look like what you expected? Use the matching problem in the classroom to assure you fully understand what was computed here.

```
In [33]: plt.hist(p_diffs)
plt.axvline(x=np.mean(p_diffs), color = 'r');
```



j. What proportion of the `p_diffs` are greater than the actual difference observed in `ab_data.csv`?

```
In [34]: obs_diff = df2.query('group=="treatment"').converted.mean() - df2.query('group=="control"').converted.mean()
p_val = (p_diffs > obs_diff).mean()
p_val
```

```
Out[34]: 0.9022
```

k. Please explain using the vocabulary you've learned in this course what you just computed in part j. What is this value called in scientific studies? What does this value mean in terms of whether or not there is a difference between the new and old pages?

Our p-value is 11.87%. Since we are using a 95% confidence interval our alpha is 5%. This is the probability of rejecting the null hypothesis when its true, also known as type 1 error.

Our p-value is greater than our alpha so we would fail to reject the null hypothesis. This means there is no statistically significant conversion rate difference between the old and new page.

l. We could also use a built-in to achieve similar results. Though using the built-in might be easier to code, the above portions are a walkthrough of the ideas that are critical to correctly thinking about statistical significance. Fill in the below to calculate the number of conversions for each page, as well as the number of individuals who received each page. Let `n_old` and `n_new` refer the number of rows associated with the old page and new pages, respectively.

```
In [35]: import statsmodels.api as sm
convert_old = df2.query('landing_page=="old_page" and converted==1')['converted'].count()
convert_new = df2.query('landing_page=="new_page" and converted==1')['converted'].count()
n_old = df2.query('landing_page=="old_page"')['converted'].count()
n_new = df2.query('landing_page=="new_page"')['converted'].count()
```

m. Now use `stats.proportions_ztest` to compute your test statistic and p-value. Here is a helpful link on using the built in.

```
In [36]: from statsmodels.stats.proportion import proportions_ztest
# Column: proportion,1 is convert_old,convert_new, [n_old,n_new], alternative="t",
stat, pval
```

```
Out[36]: (1.3109241984234394, 0.3905098312759245)
```

n. What do the z-score and p-value you computed in the previous question mean for the conversion rates of the old and new pages? Do they agree with the findings in parts j, and k?

The Z score tells us that our data is at 1.31 standard deviations away from the mean. A p value of 0.905 is way above our alpha of 0.05. Therefore suggesting the data is not statistically significant so we fail to reject the null.

Since we fail to reject the null, this is in agreement with the answers in parts j and k

### Part III - A regression approach

1. In this final part you will see that the result you achieved in the A/B test in Part II above can also be achieved by performing regression.

a. Since each row is either a conversion or no conversion, what type of regression should you be performing in this case?

Logistic Regression

b. The goal is to use `statsmodels` to fit the regression model you specified in part a. to see if there is a significant difference in conversion based on which page a customer receives. However, you first need to create in `df2` a column for the intercept, and create a dummy variable column for which page each user received. Add an intercept column, as well as an `ab_page` column, which is 1 when an individual receives the `treatment` and 0 if `control`.

```
In [37]: df2['intercept'] = 1
df2['ab_page'] = pd.get_dummies(df2['group'])['treatment']
df2.head()
```

```
Out[37]:
```

|   | user_id | timestamp                  | group     | landing_page | converted | intercept | ab_page |
|---|---------|----------------------------|-----------|--------------|-----------|-----------|---------|
| 0 | 851104  | 2017-01-21 22:11:48.556739 | control   | old_page     | 0         | 1         | 0       |
| 1 | 804228  | 2017-01-12 08:01:45.159739 | control   | old_page     | 0         | 1         | 0       |
| 2 | 661590  | 2017-01-11 16:55:06.154213 | treatment | new_page     | 0         | 1         | 1       |
| 3 | 853541  | 2017-01-08 18:28:03.143765 | treatment | new_page     | 0         | 1         | 1       |
| 4 | 864975  | 2017-01-21 01:52:26.210827 | control   | old_page     | 1         | 1         | 0       |

c. Use `statsmodels` to instantiate your regression model on the two columns you created in part b., then fit the model using the two columns you created in part b. to predict whether or not an individual converts.

```
In [38]: logit_mod = sm.Logit(df2['converted'], df2[['intercept', 'ab_page']])
results = logit_mod.fit()
Optimization terminated successfully.
Iteration function value: 0.366118
Iterations: 6
```

d. Provide the summary of your model below, and use it as necessary to answer the following questions.

```
In [39]: results.summary2()
```

```
Out[39]:
```

|                     | Model:           | Logit           | Pseudo R-squared: | 0.000 |
|---------------------|------------------|-----------------|-------------------|-------|
| Dependent Variable: | converted        | AIC:            | 212780.3502       |       |
| Date:               | 2021-04-06 01:32 | BIC:            | 212801.5095       |       |
| No. Observations:   | 290584           | Log-Likelihood: | -1.0639e+05       |       |
| Df Model:           | 1                | LL-Null:        | -1.0639e+05       |       |
| Df Residuals:       | 290582           | LLR p-value:    | 0.18988           |       |
| Converged:          | 1.0000           | Scale:          | 1.0000            |       |
| No. Iterations:     | 6.0000           |                 |                   |       |

```
Coef. Std.Err. z P>|z| [0.025 0.975]
intercept -1.9888 0.0081 -246.6690 0.0000 -2.0046 -1.9731
ab_page -0.0150 0.0114 -1.3109 0.1899 -0.0374 0.0074
```

e. What is the p-value associated with `ab_page`? Why does it differ from the value you found in Part II?

Hint: What are the null and alternative hypotheses associated with your regression model, and how do they compare to the null and alternative hypotheses in Part II?

The p-value is 0.1899.

the null hypothesis from the regression model states that the parameter is equal to zero, the alternate states that it does not equal zero which makes it a 2 sided hypothesis test.

the null and alternate hypothesis in part II are one sided tests. therefore the resulting p value from the regression is double the answer in part II

f. Now, you are considering other factors that might influence whether or not an individual converts. Discuss why it is a good idea to consider other factors to add into your regression model. Are there any disadvantages to adding additional terms into your regression model?

another variable to consider is the timestamp, it would be a good to add to the model to see if it can increase the accuracy of the model for predicting if a user will convert or not. the disadvantage of this is that it could introduce multicollinearity which needs to be tested if it should be removed, and it can introduce the ideas of change aversion and the novelty effect of the newly designed website

```
In [40]: results.params
```

```
Out[40]:
```

|         | intercept | -1.988777 |
|---------|-----------|-----------|
| ab_page | -0.014989 |           |
| dtype:  | float64   |           |

```
In [41]: 1/np.exp(_)
Out[41]:
```

|         | intercept | 7.306593 |
|---------|-----------|----------|
| ab_page | 0.985169  |          |
| dtype:  | float64   |          |

interpretation: for each 1 unit decrease in `ab_page`, conversion is 1.015 times as likely holding all else constant. further interpreted as: if `ab_page` is 0, which is the old control page, then the chance of conversion goes up by 1.015 times.

g. Now along with testing if the conversion rate changes for different pages, also add an effect based on which country a user lives in. You will need to read in the `countries.csv` dataset and merge together your datasets on the appropriate rows. Here are the docs for joining tables.

Does it appear that country had an impact on conversions? Don't forget to create dummy variables for these country columns - Hint: You will need two columns for the three dummy variables. Provide the statistical output as well as a written response for this question.

```
In [42]: # read in the countries CSV file
# create a new dataframe and name it df3 that joins df2 and the countries CSV file.
countries = pd.read_csv('countries.csv')
df3 = df2.join(countries.set_index('user_id'), on='user_id')
df3.head()
```

```
Out[42]:
```

|   | user_id | timestamp                  | group     | landing_page | converted | intercept | ab_page | country |
|---|---------|----------------------------|-----------|--------------|-----------|-----------|---------|---------|
| 0 | 851104  | 2017-01-21 22:11:48.556739 | control   | old_page     | 0         | 1         | 0       | US      |
| 1 | 804228  | 2017-01-12 08:01:45.159739 | control   | old_page     | 0         | 1         | 0       | US      |
| 2 | 661590  | 2017-01-11 16:55:06.154213 | treatment | new_page     | 0         | 1         | 1       | US      |
| 3 | 853541  | 2017-01-08 18:28:03.143765 | treatment | new_page     | 0         | 1         | 1       | US      |
| 4 | 864975  | 2017-01-21 01:52:26.210827 | control   | old_page     | 1         | 1         | 0       | US      |

```
In [43]: # get the dummy values for the countries read in and give them descriptive column names
df3[['CA', 'UK', 'US']] = pd.get_dummies(df3['country'])
df3.tail()
```

```
Out[43]:
```

|        | user_id | timestamp     | group   | landing_page | converted | intercept | ab_page | country | CA | UK | US |
|--------|---------|---------------|---------|--------------|-----------|-----------|---------|---------|----|----|----|
| 244473 | 751197  | 2017-01-03    | control | old_page     | 0         | 1         | 0       | US      | 0  | 0  | 0  |
| 244474 | 945152  | 005157.078372 | control | old_page     | 0         | 1         | 0       | US      | 0  | 0  | 0  |
| 244475 | 734608  | 2017-01-22    | control | old_page     | 0         | 1         |         |         |    |    |    |



|           | No. Iterations: |          | 6.0000    |        |         |         |  |
|-----------|-----------------|----------|-----------|--------|---------|---------|--|
|           | Coef.           | Std.Err. | z         | P> z   | [0.025  | 0.975]  |  |
| Intercept | -1.9719         | 0.0163   | -120.9419 | 0.0000 | -2.0039 | -1.9400 |  |
| ab_page   | -0.0205         | 0.0137   | -1.4998   | 0.1337 | -0.0473 | 0.0063  |  |
| UK        | -0.0174         | 0.0377   | -0.4617   | 0.6443 | -0.0913 | 0.0565  |  |
| CA        | -0.0056         | 0.0188   | -0.2965   | 0.7668 | -0.0424 | 0.0313  |  |
| CA_page   | -0.0470         | 0.0538   | -0.8743   | 0.3819 | -0.1524 | 0.0584  |  |
| UK_page   | 0.0314          | 0.0266   | 1.1795    | 0.2382 | -0.0208 | 0.0835  |  |
| Tuesday   | -0.0153         | 0.0202   | -0.7551   | 0.4502 | -0.0549 | 0.0243  |  |
| Wednesday | -0.0076         | 0.0211   | -0.3604   | 0.7186 | -0.0490 | 0.0338  |  |
| Thursday  | -0.0110         | 0.0212   | -0.5212   | 0.6022 | -0.0525 | 0.0305  |  |
| Friday    | -0.0425         | 0.0212   | -1.9995   | 0.0456 | -0.0841 | -0.0008 |  |
| Saturday  | -0.0030         | 0.0211   | -0.1437   | 0.8857 | -0.0443 | 0.0383  |  |
| Sunday    | -0.0254         | 0.0211   | -1.2026   | 0.2291 | -0.0668 | 0.0160  |  |

In [62]: # Get coefficients of logistic model results.params

Out[62]: Intercept -1.971926  
ab\_page -0.020491  
CA -0.017398  
UK -0.005572  
CA\_page -0.047021  
UK\_page 0.031366  
Tuesday -0.015257  
Wednesday -0.007607  
Thursday -0.011025  
Friday -0.042489  
Saturday -0.003028  
Sunday -0.025414  
dtype: float64

In [63]: # Interpret coefficients greater than 1 np.exp(results.params)

Out[63]: ab\_page 0.139188  
Intercept 0.979718  
CA 0.940725  
CA\_page 0.954752  
Friday 0.958752  
Saturday 0.997271  
Sunday 0.975001  
Thursday 0.989151  
Tuesday 0.985001  
UK 0.985572  
UK\_page 1.032366  
Wednesday 0.993607

```
In [62]: # get coefficients of logistic model
results.params
```

```
Out[62]: Intercept      -1.971926
          ab_page       -0.020491
          CA            -0.017398
          UK            -0.005572
          CA_page       -0.047021
          UK_page       -0.031366
          Tuesday       -0.015257
          Wednesday     -0.007607
          Thursday      -0.011035
          Friday        -0.042489
          Saturday      -0.003028
          Sunday        -0.025414
          dtype: float64
```

```
In [63]: # interpret coefficients greater than 1
np.exp(results.params)
```

```
Out[63]: Intercept      0.139188
          ab_page       0.979718
          CA            0.979753
          UK            0.994444
          CA_page       0.954067
          UK_page       1.031863
          Tuesday       0.984858
          Wednesday     0.992422
          Thursday      0.989026
          Friday        0.959401
          Saturday      0.996976
          Sunday        0.974906
          dtype: float64
```

```
In [64]: # interpret coefficients below 1
1/np.exp(results.params)
```

```
Out[64]: Intercept      7.184503
          ab_page       1.020702
          CA            1.017550
          UK            1.005587
          CA_page       1.048144
          UK_page       0.969121
          Tuesday       1.015374
          Wednesday     1.007636
          Thursday      1.011096
          Friday        1.043404
          Saturday      1.003033
          Sunday        1.025740
          dtype: float64
```

From these coefficients it seems that the day of the week has no significant impact on conversion rates.

At best relative to monday, students using the website on Friday are more liekly to convert by a factor of 1.043.

This seems like a practically insignificant difference.

All p values in this model are above 0.05 so we fail to reject the null hypothesis

However this experiment has only been run for 21 days so there just might not be enough data.

I would recoomend to run the experiment for a longer time

## Extra models are created below to test if accuracy increases by removing certain variables

```
In [65]: # build identical model however remove interaction
logit_mod = sm.Logit(df4['converted'], df4[['Intercept', 'ab_page', 'CA', 'UK', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']])
results = logit_mod.fit()
results.summary2()
```

Optimization terminated successfully.  
Current function value: 0.366103  
Iterations: 6

```
Out[65]:
```

| Model:              | Logit            | Pseudo R-squared: | AIC: | 0.000       |
|---------------------|------------------|-------------------|------|-------------|
| Dependent Variable: | converted        |                   | AIC: | 212787.5901 |
| Date:               | 2021-04-06 01:33 |                   | BIC: | 212893.3866 |

|                   |        |              |             |
|-------------------|--------|--------------|-------------|
| No. Observations: | 290584 |              |             |
| Df Model:         | 9      | LL-Null:     | -1.0639e+05 |
| Df Residuals:     | 290574 | LLR p-value: | 0.31314     |
| Converged:        | 1.0000 | Scale:       | 1.0000      |
| No. Iterations:   | 6.0000 |              |             |

|           | Coef.   | Std.Err. | z         | P> z   | [0.025  | 0.975]  |
|-----------|---------|----------|-----------|--------|---------|---------|
| Intercept | -1.9747 | 0.0159   | -124.2803 | 0.0000 | -2.0058 | -1.9436 |
| ab_page   | -0.0149 | 0.0114   | -1.3018   | 0.1930 | -0.0373 | 0.0075  |
| CA        | -0.0407 | 0.0269   | -1.5138   | 0.1301 | -0.0934 | 0.0120  |
| UK        | 0.0100  | 0.0133   | 0.7551    | 0.4502 | -0.0160 | 0.0361  |
| Tuesday   | -0.0153 | 0.0202   | -0.7558   | 0.4498 | -0.0549 | 0.0243  |
| Wednesday | -0.0076 | 0.0211   | -0.3594   | 0.7193 | -0.0490 | 0.0338  |
| Thursday  | -0.0111 | 0.0212   | -0.5239   | 0.6004 | -0.0526 | 0.0304  |
| Friday    | -0.0425 | 0.0212   | -2.0002   | 0.0455 | -0.0842 | -0.0009 |
| Saturday  | -0.0031 | 0.0211   | -0.1457   | 0.8841 | -0.0444 | 0.0382  |
| Sunday    | -0.0254 | 0.0211   | -1.2021   | 0.2293 | -0.0668 | 0.0160  |

```
In [66]: # get coefficients of logistic model
results.params
```

```
Out[66]: Intercept      -1.974700
          ab_page       -0.014885
          CA            -0.040695
          UK            0.010040
          Tuesday       -0.015270
          Wednesday     -0.007586
          Thursday      -0.011091
          Friday        -0.042502
          Saturday      -0.003070
          Sunday        -0.025405
          dtype: float64
```

```
In [67]: # interpret coefficients greater than 1
np.exp(results.params)
```

```
Out[67]: Intercept      0.138803
          ab_page       0.985225
          CA            0.960122
          UK            1.010090
          Tuesday       0.984846
          Wednesday     0.992443
          Thursday      0.988970
          Friday        0.958388
          Saturday      0.996934
          Sunday        0.974915
          dtype: float64
```

```
In [68]: # interpret coefficients below 1
1/np.exp(results.params)
```

```
Out[68]: Intercept      7.204456
          ab_page       1.014397
          CA            1.041535
          UK            0.990010
          Tuesday       1.015387
          Wednesday     1.007615
          Thursday      1.011153
          Friday        1.043419
          Saturday      1.003075
          Sunday        1.025730
          dtype: float64
```

```
In [69]: # build similar model further simplifying by removing countries
logit_mod = sm.Logit(df4['converted'], df4[['Intercept', 'ab_page', 'Tuesday', 'Wednesday', 'Thursday', 'Friday', 'Saturday', 'Sunday']])
results = logit_mod.fit()
results.summary2()
```

Optimization terminated successfully.  
Current function value: 0.366109  
Iterations: 6

```
Out[69]:
```

| Model:              | Logit            | Pseudo R-squared: | AIC: | 0.000       |
|---------------------|------------------|-------------------|------|-------------|
| Dependent Variable: | converted        |                   | AIC: | 212786.8299 |
| Date:               | 2021-04-06 01:33 |                   | BIC: | 212871.4671 |

|                   |        |                 |             |
|-------------------|--------|-----------------|-------------|
| No. Observations: | 290584 |                 |             |
| Df Model:         | 7      | Log-Likelihood: | -1.0639e+05 |
| Df Residuals:     | 290576 | LLR p-value:    | 0.40444     |
| Converged:        | 1.0000 | Scale:          | 1.0000      |
| No. Iterations:   | 6.0000 |                 |             |

|           | Coef.   | Std.Err. | z         | P> z   | [0.025  | 0.975]  |
|-----------|---------|----------|-----------|--------|---------|---------|
| Intercept | -1.9742 | 0.0155   | -127.7522 | 0.0000 | -2.0045 | -1.9439 |
| ab_page   | -0.0149 | 0.0114   | -1.3058   | 0.1916 | -0.0373 | 0.0075  |
| Tuesday   | -0.0152 | 0.0202   | -0.7523   | 0.4519 | -0.0548 | 0.0244  |
| Wednesday | -0.0076 | 0.0211   | -0.3620   | 0.7173 | -0.0490 | 0.0337  |
| Thursday  | -0.0111 | 0.0212   | -0.5264   | 0.5986 | -0.0526 | 0.0303  |
| Friday    | -0.0425 | 0.0212   | -1.9991   | 0.0456 | -0.0841 | -0.0008 |
| Saturday  | -0.0030 | 0.0211   | -0.1446   | 0.8850 | -0.0443 | 0.0383  |
| Sunday    | -0.0253 | 0.0211   | -1.1986   | 0.2307 | -0.0667 | 0.0161  |

```
In [70]: # get coefficients of logistic model
results.params
```

```
Out[70]: Intercept      -1.974176
          ab_page       -0.014931
          Tuesday       -0.015200
          Wednesday     -0.007641
          Thursday      -0.011145
          Friday        -0.042479
          Saturday      -0.003047
          Sunday        -0.025329
          dtype: float64
```

```
In [71]: # interpret coefficients greater than 1
np.exp(results.params)
```

```
Out[71]: Intercept      0.138876
          ab_page       0.985180
          Tuesday       0.984915
          Wednesday     0.992388
          Thursday      0.988917
          Friday        0.958411
          Saturday      0.996958
          Sunday        0.974989
          dtype: float64
```

```
In [72]: # interpret coefficients below 1
1/np.exp(results.params)
```

```
Out[72]: Intercept      7.200685
          ab_page       1.015043
          Tuesday       1.015136
          Wednesday     1.007670
          Thursday      1.011207
          Friday        1.043394
          Saturday      1.003051
          Sunday        1.025652
          dtype: float64
```

We can see removing the countries and interactions has no effect on conversion rates for day of week as the coefficients remain unchanged.

Final recommendation is to run experiment for a longer period of time. As for the data collected here, we have failed to reject the null and see no statistically significant difference between the old and new websites when it comes to converting users

```
In [73]: from subprocess import call
         call(['python', '-m', 'nbconvert', 'Analyze_ab_test_results_notebook.ipynb'])
```

```
Out[73]: 1
```